Week 11

- Photo OCR:
  a. Problem Description and Pipeline:
  Photo OCR: Photo Optical Character Recognition
  First, given a picture it has to look through the image and detect where there is text in the picture.

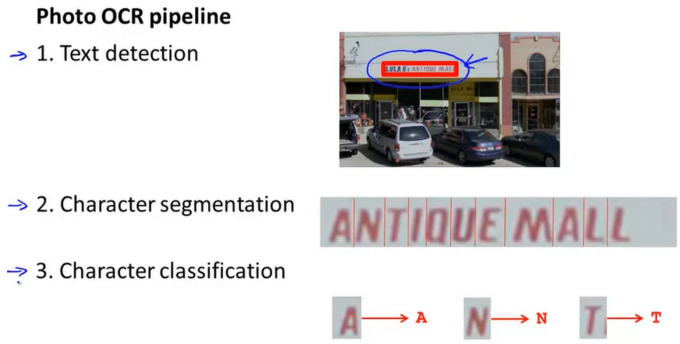

  b. Sliding Windows:
  these rectangles that you're trying to find can have different aspect

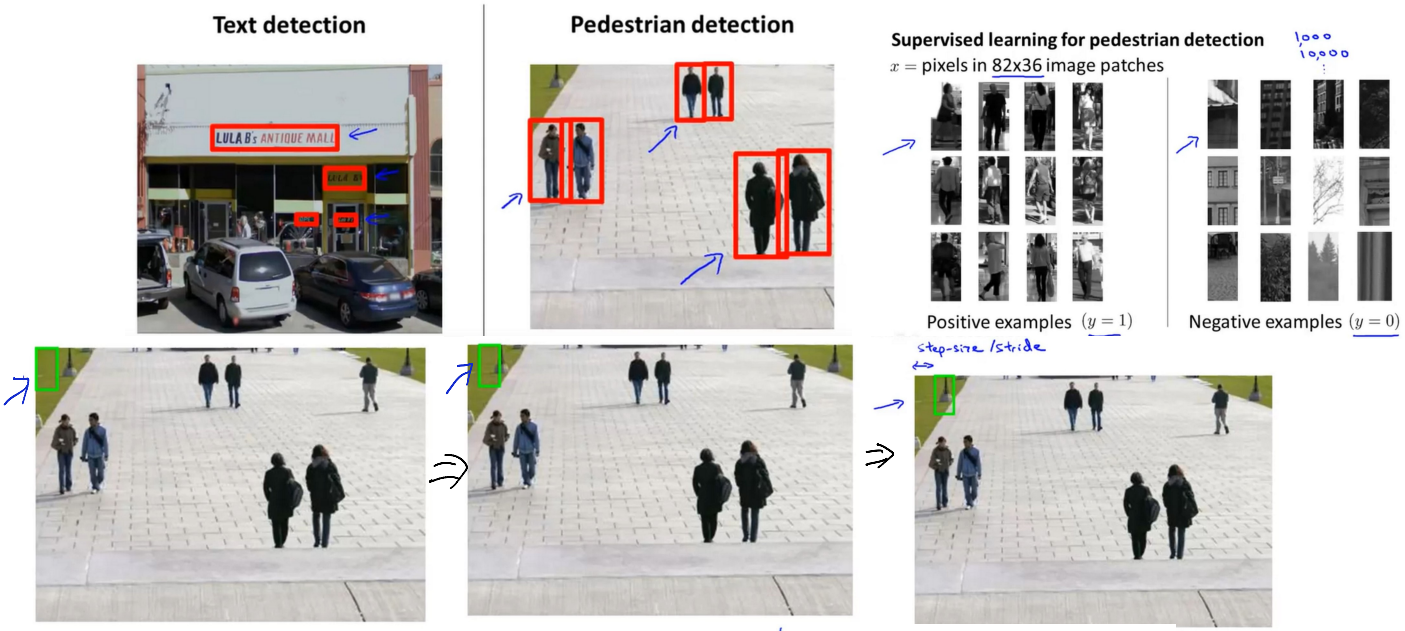

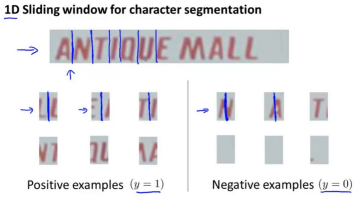

c. Getting Lots of Data and Artificial Data:
Artificial data synthesis

{ create from begining.
already have a small label training set and we amplify training set.

Synthesizing data by introducing distortions

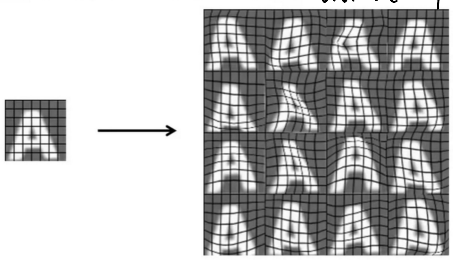

→ Distortion introduced should be representation of the type of noise/distortions in the test set.

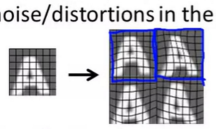

→ Audio:
Background noise,
bad cellphone connection

→ Usually does not help to add purely random/meaningless noise to your data.

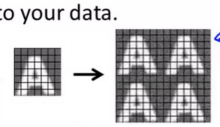

→ x_i = intensity (brightness) of pixel i
→ $x_i \leftarrow x_i +$ random noise

Discussion on getting more data

1. Make sure you have a low bias classifier before expending the effort. (Plot learning curves). E.g. keep increasing the number of features/number of hidden units in neural network until you have a low bias classifier.
2. "How much work would it be to get 10x as much data as we currently have?"
   - Artificial data synthesis
   - Collect/label it yourself
   - "Crowd source" (E.g. Amazon Mechanical Turk)

d. Ceiling Analysis What Part of the Pipeline to Work on Next:
Tell you what parts of the pipeline might be the best use of your time to work on.

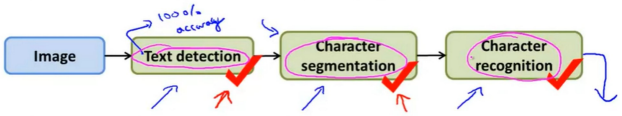

What part of the pipeline should you spend the most time trying to improve?

Component	Accuracy	overall system.
Overall system	72%	↓ 17%
→ Text detection	89%	↓ 1%
Character segmentation	90%	↓ 10%
Character recognition	100%	

how much could you possibly gain if one of these components became absolutely perfect?